

Fashion-MNIST

Exploring The Dataset

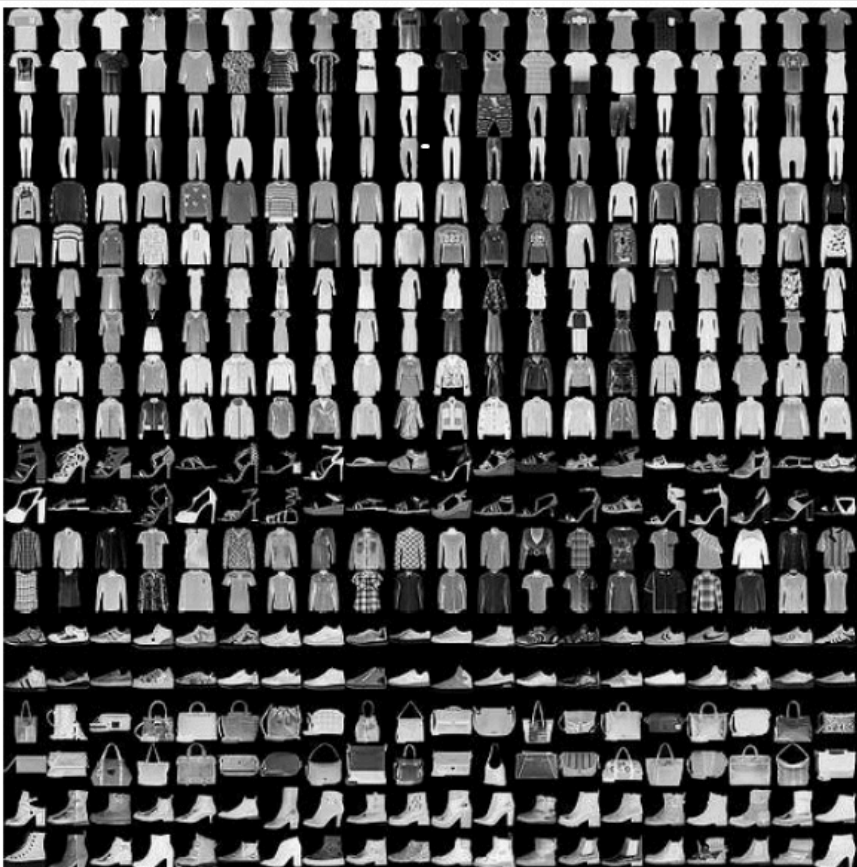
Import Libraries

```
In [1]: import numpy as np
        from keras.models import Sequential
        from keras.layers import Dropout, Dense, Flatten
        from tensorflow.keras.optimizers import SGD, Adam
        from keras.layers import Conv2D, MaxPooling2D
        from tensorflow.keras.utils import to_categorical
        from keras.datasets import fashion_mnist
        from matplotlib import pyplot as plt
        import warnings
        warnings.filterwarnings("ignore")
        import seaborn as sns
```

Load Data


```
In [2]: #Lets start by loading the Cifar10 data
        (x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()
        print(x_train.shape, y_train.shape)
        print(x_test.shape, y_test.shape)

(60000, 28, 28) (60000,)
(10000, 28, 28) (10000,)
```

Label	Description	Examples
0	T-Shirt/Top	
1	Trouser	
2	Pullover	
3	Dress	
4	Coat	
5	Sandals	
6	Shirt	
7	Sneaker	
8	Bag	
9	Ankle boots	

In [3]: `plt.imshow(x_train[4])`

Out[3]: `<matplotlib.image.AxesImage at 0x205e7b9c460>`



Process Data

```
In [4]: x_train = x_train.astype('float32')/255.0
        x_test = x_test.astype('float32')/255.0
```

```
In [5]: y_train = to_categorical(y_train, 10)
        y_test = to_categorical(y_test, 10)
```

```
In [6]: print (y_train.shape)

(60000, 10)
```

Define Model

```
In [7]: model = Sequential()
        model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=(28, 28, 1), padding='same'))
        model.add(Conv2D(64, (3, 3), activation='relu', padding='same'))
        model.add(MaxPooling2D(pool_size=(2, 2), padding='same'))
        model.add(Dropout(0.3))
        model.add(Flatten())
        model.add(Dense(256, activation='relu'))
        model.add(Dropout(0.5))
        model.add(Dense(10, activation='softmax'))
```

Model Summary

```
In [8]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Pa
conv2d (Conv2D)	(None, 28, 28, 32)	
conv2d_1 (Conv2D)	(None, 28, 28, 64)	
max_pooling2d (MaxPooling2D)	(None, 14, 14, 64)	
dropout (Dropout)	(None, 14, 14, 64)	
flatten (Flatten)	(None, 12544)	
dense (Dense)	(None, 256)	3,232,906
dropout_1 (Dropout)	(None, 256)	
dense_1 (Dense)	(None, 10)	

Total params: 3,232,906 (12.33 MB)


Trainable params: 3,232,906 (12.33 MB)


Compile Model


```
In [9]: model.compile(loss='categorical_crossentropy', optimizer=Adam(), metrics=['accuracy'])
```


Fit data to model


```
In [10]: epochs = 20
batch_size = 128
hist = model.fit(x_train, y_train, batch_size=batch_size,
                 epochs=epochs, verbose=1,
                 validation_data=(x_test, y_test))
print("The model has successfully trained")
```


Epoch 1/20
469/469  56s 117ms/step - accuracy: 0.7780 - loss: 0.6247 - val_accuracy: 0.8886 - val_loss: 0.3040


Epoch 2/20
469/469  59s 126ms/step - accuracy: 0.8901 - loss: 0.3072 - val_accuracy: 0.9010 - val_loss: 0.2646


Epoch 3/20
469/469  63s 134ms/step - accuracy: 0.9094 - loss: 0.2481 - val_accuracy: 0.9095 - val_loss: 0.2378


Epoch 4/20
469/469  69s 146ms/step - accuracy: 0.9193 - loss: 0.2163 - val_accuracy: 0.9083 - val_loss: 0.2488


Epoch 5/20
469/469  73s 155ms/step - accuracy: 0.9271 - loss: 0.1946 - val_accuracy: 0.9166 - val_loss: 0.2222


Epoch 6/20
469/469  66s 140ms/step - accuracy: 0.9375 - loss: 0.1655 - val_accuracy: 0.9242 - val_loss: 0.2082


Epoch 7/20
469/469  66s 140ms/step - accuracy: 0.9422 - loss: 0.1574 - val_accuracy: 0.9245 - val_loss: 0.2128


Epoch 8/20
469/469  62s 133ms/step - accuracy: 0.9489 - loss: 0.1374 - val_accuracy: 0.9283 - val_loss: 0.2076


Epoch 9/20
469/469  63s 135ms/step - accuracy: 0.9517 - loss: 0.1255 - val_accuracy: 0.9295 - val_loss: 0.2075


Epoch 10/20
469/469  66s 140ms/step - accuracy: 0.9561 - loss: 0.1163 - val_accuracy: 0.9343 - val_loss: 0.2020


Epoch 11/20
469/469  65s 138ms/step - accuracy: 0.9613 - loss: 0.1027 - val_accuracy: 0.9327 - val_loss: 0.2145


Epoch 12/20
469/469  62s 132ms/step - accuracy: 0.9638 - loss: 0.0969 - val_accuracy: 0.9282 - val_loss: 0.2274


Epoch 13/20
469/469  58s 123ms/step - accuracy: 0.9664 - loss: 0.0874 - val_accuracy: 0.9323 - val_loss: 0.2220


Epoch 14/20
469/469  59s 126ms/step - accuracy: 0.9702 - loss: 0.0772 - val_accuracy: 0.9280 - val_loss: 0.2378


Epoch 15/20
469/469  66s 142ms/step - accuracy: 0.9710 - loss: 0.0749 - val_accuracy: 0.9355 - val_loss: 0.2297

Epoch 16/20
469/469  66s 141ms/step - accuracy: 0.9745 - loss: 0.0673 - val_accuracy: 0.9328 - val_loss: 0.2449

Epoch 17/20
469/469  63s 135ms/step - accuracy: 0.9756 - loss: 0.0648 - val_accuracy: 0.9331 - val_loss: 0.2521

Epoch 18/20
469/469  64s 136ms/step - accuracy: 0.9785 - loss: 0.0576 - val_accuracy: 0.9327 - val_loss: 0.2535

Epoch 19/20
469/469  63s 133ms/step - accuracy: 0.9793 - loss: 0.0564 - val_accuracy: 0.9335 - val_loss: 0.2568

Epoch 20/20
469/469  58s 124ms/step - accuracy: 0.9803 - loss: 0.0543 - val_accuracy: 0.9335 - val_loss: 0.2568

ccuracy: 0.9376 - val_loss: 0.2810
The model has successfully trained

Save the Model to use for later

```
In [11]: model.save_weights('FashionMNIST.weights.h5')  
print("Saving the model as FashionMNIST.h5")
```

Saving the model as FashionMNIST.h5

Evaluate Model

```
In [12]: score = model.evaluate(x_test, y_test, verbose=0)  
print('Test loss:', score[0])  
print('Test accuracy:', score[1])
```

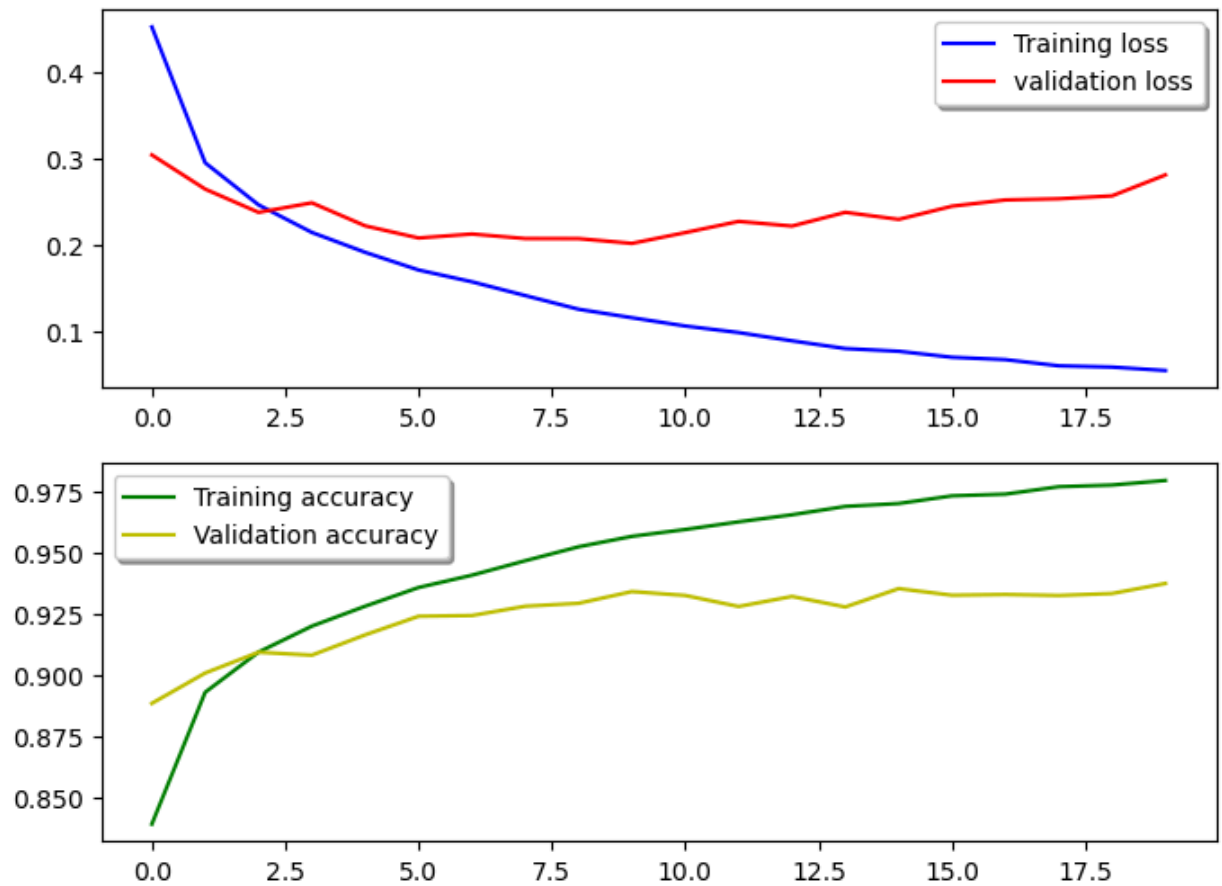
Test loss: 0.27781593799591064

Test accuracy: 0.9376000165939331

This model predicted 93.76% of Test Images correctly, which indicates that the model did pretty good job in generalizing the data.

Plotting the Training and Validation Curves

```
In [13]: fig, ax = plt.subplots(2,1, figsize=(8, 6))  
ax[0].plot(hist.history['loss'], color='b', label="Training loss")  
ax[0].plot(hist.history['val_loss'], color='r',  
            label="validation loss", axes = ax[0])  
legend = ax[0].legend(loc='best', shadow=True)  
  
ax[1].plot(hist.history['accuracy'], color='g', label="Training accuracy")  
ax[1].plot(hist.history['val_accuracy'], color='y',  
            label="Validation accuracy")  
legend = ax[1].legend(loc='best', shadow=True)
```

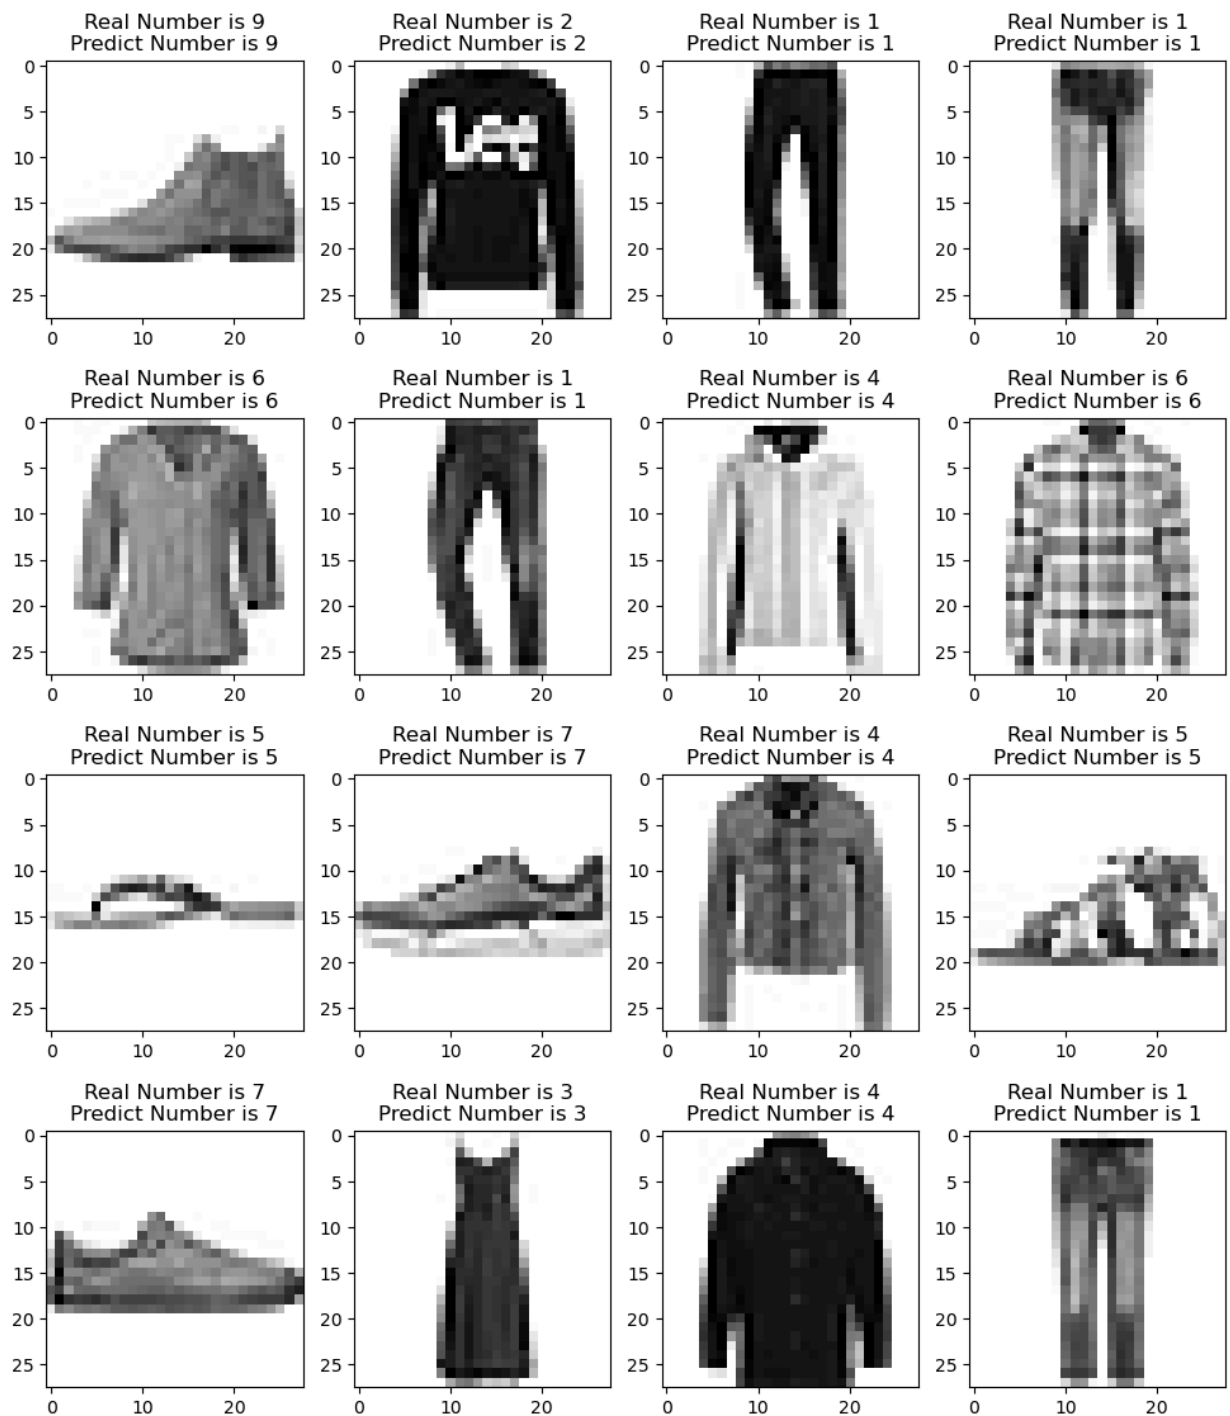


Make Predictions

```
In [14]: y_pred = model.predict(x_test)
X_test__ = x_test.reshape(x_test.shape[0], 28, 28)

fig, axis = plt.subplots(4, 4, figsize=(12, 14))
for i, ax in enumerate(axis.flat):
    ax.imshow(X_test__[i], cmap='binary')
    ax.set(title = f"Real Number is {y_test[i].argmax()}\nPredict Number is {y_pred[i]})")
```


313/313 ————— 3s 9ms/step



Testing Model with unknown image

```
In [15]: import tensorflow as tf
import cv2 as cv
from matplotlib import pyplot as plt
img=cv.imread('sample_image.jpg')
plt.imshow(img)
```

```
Out[15]: <matplotlib.image.AxesImage at 0x205802764a0>
```

```
In [16]: img.shape
```

```
Out[16]: (740, 740, 3)
```

```
In [17]: img = cv.cvtColor(img, cv.COLOR_BGR2GRAY)
img = cv.resize(img, (28, 28))
```

```
In [18]: img.shape
```

```
Out[18]: (28, 28)
```

```
In [19]: from tensorflow.keras.utils import img_to_array
from tensorflow import keras
model = keras.models.load_model('FashionMNIST.h5')
```

WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile_metrics` will be empty until you train or evaluate the model.

```
In [20]: model.input_shape
```

```
Out[20]: (None, 28, 28, 1)
```

```
In [21]: import numpy as np
from tensorflow.keras.utils import img_to_array
x = img_to_array(img)
x = np.expand_dims(x, axis=0)
preds = np.argmax(model.predict(x))
print(preds)
```


1/1 ————— 0s 219ms/step

2

Incorrectly Predicted Class

```
In [22]: img2=cv.imread('sample_image2.jpg')  
plt.imshow(img2)
```

```
Out[22]: <matplotlib.image.AxesImage at 0x205e76ae320>
```



```
In [23]: img2.shape
```

```
Out[23]: (240, 210, 3)
```

```
In [24]: img2 = cv.cvtColor(img2, cv.COLOR_BGR2GRAY)  
img2 = cv.resize(img2, (28, 28))
```

```
In [25]: img2.shape
```

```
Out[25]: (28, 28)
```

```
In [26]: import numpy as np  
from tensorflow.keras.utils import img_to_array  
x = img_to_array(img2)  
x = np.expand_dims(x, axis=0)  
preds = np.argmax(model.predict(x))  
print(preds)
```

```
1/1 ————— 0s 31ms/step  
2
```

Classification Report

```
In [27]: from sklearn.metrics import confusion_matrix, classification_report
y_pred_classes = np.argmax(y_pred,axis = 1)
y_true = np.argmax(y_test,axis = 1)
classes = ['T-shirt/Top','Trouser','Pullover','Dress','Coat','Sandal','Shirt',
           'Sneaker','Bag','Ankle Boot']
print(classification_report(y_true, y_pred_classes, target_names = classes))
```

	precision	recall	f1-score	support
T-shirt/Top	0.91	0.89	0.90	1000
Trouser	1.00	0.98	0.99	1000
Pullover	0.90	0.89	0.90	1000
Dress	0.94	0.94	0.94	1000
Coat	0.89	0.91	0.90	1000
Sandal	0.99	0.99	0.99	1000
Shirt	0.81	0.82	0.82	1000
Sneaker	0.97	0.98	0.97	1000
Bag	0.99	0.99	0.99	1000
Ankle Boot	0.98	0.97	0.98	1000
accuracy			0.94	10000
macro avg	0.94	0.94	0.94	10000
weighted avg	0.94	0.94	0.94	10000

Confusion Matrix

```
In [28]: confusion_matrix(y_true, y_pred_classes)
sns.heatmap(confusion_matrix(y_true, y_pred_classes),annot=True,cmap='Set3',
            linewidths=0.3, linecolor='black',fmt='.0f')
```

Out[28]: <Axes: >

